



Under the Hood of the LLM Revolution



Three questions the terminal is about to answer

Gourav Shah · School of DevOps & AI · Deep Dive (Part 2)

Softmax turns two raw scores into weights that sum to one




the exact same squashing function that turns raw logits into probabilities

the blend isn't hand-tuned — it falls straight out of this arithmetic

The exact function that also drives sampling later



"Close together" means a small angle, not a short distance



dividing by each vector's length is what makes this
about direction, not length


Dividing by length is what makes this about direction

Temperature reshapes the whole distribution before sampling



Divides every raw score before softmax runs

Top-p trims which candidates are even eligible



A different mechanism from temperature

Now: run the sampling script, watch the distribution move

- 1 sampling_demo.py — watch the ten samples spread as temperature climbs
- 2 same script — watch top-p exclude the tail

labs/m5 — live now

Same three ideas — proven live, not just drawn