


Two questions this deep dive answers



Proof, not a spot-check

Gourav Shah · School of DevOps & AI · Deep Dive (Part 2)

A unit test checks the type; an eval checks the answer

a unit test

right type of thing,
doesn't crash


an eval

the right answer, or close
enough to it

a model can never error and still be quietly, systematically wrong


Never erroring isn't the same as being right

50 held-out houses in, one pass/fail verdict out



The exact idea behind Module 3's own scoring step

Swap two fields, and every request still returns 200 OK



nothing about this failure is visible unless something checks the **quality** of the answer

Nothing crashes, nothing looks wrong outside

Three signals worth watching once it's live

- 1 latency — how long a /predict call takes, end to end
- 2 error rate — the share of requests that fail outright
- 3 prediction drift — the output distribution shifting

an eval catches a bug before release; monitoring catches a problem that shows up later, from real traffic an eval never saw

A different window in time than the eval above

HAND-OFF TO THE TERMINAL

Cost is requests $\times$ tokens $\times$ price

2,000 requests/day $\times$ 800 in + 300 out
tokens $\times$ \$0.50 / \$1.50 per M tokens

$\approx$ \$51/month

the same feature at 200,000
requests/day — nothing else changes

$\approx$ \$5,100/month

Any one of the three can move on its own

Now run `eval_harness.py`, then `eval_harness.py --simulate-regression`