


Two things this deep dive goes under the hood on



Proof, not a spot-check

Gourav Shah · School of DevOps & AI · Deep Dive (Part 2)

Grading a multiple-choice exam is a test suite; grading an essay is an eval suite

multiple-choice exam

one correct bubble,
checked by a machine

no judgment involved

an essay

a rubric, a range of
acceptable answers


graded as a quality, not a match

a model can never error and still be quietly, systematically wrong


The row that trips people up

Same four stages, same order, a different artifact moving through each one

MLOps




LLMOps



different artifact, same discipline, made concrete enough to build a pipeline from


You already run this discipline — for a chain of microservice calls



"the request took 3 seconds" isn't useful; "step 2 took 2.8 of those 3" is

One user request, several downstream calls

One trace_id, five spans: reason, act, observe, reason, respond



HAND-OFF TO THE TERMINAL

Four signals a trace makes answerable that one log line never could

- 1 steps per task
- 2 tool-call error rate, per tool
- 3 cost per completed task, not per request
- 4 whether the loop actually stopped where you expect

Gourav Shah · School of DevOps & AI · Now run `pipeline_stages.py`, then `agent_trace.py`