

Frontier — What's Happening in GenAI





Sort the headlines fast. Keep building.

Gourav Shah · School of DevOps & AI · Lesson + Lab + Quiz




**Three headlines in one week, and a teammate asking
“does any of this change what we build?”**

Read all of it and you ship nothing; read none and you miss the one that matters



the habit in between
two questions, five seconds

Two ways to get this wrong




The habit: triage, don't read everything

A senior engineer never reads the whole changelog



Two fast questions, then most of it gets filed

Question one: which region of the map does this land in?



Machine learning, deep learning, the LLM, or the agent?

Question two: what does it change for you as a builder?


- 1 something you genuinely couldn't build before
- 2 the same thing — cheaper, or easier to reach
- 3 how much you check before you trust the answer
- 4 the plumbing and standards you build against

Four kinds of impact, one quick call




The five shapes GenAI news keeps taking

Every name here will look dated; the shapes will not




A snapshot, not a permanent map

A reasoning model spends more time thinking inside one call




Slower and more costly, steadier on multi-step work

One model now takes text, images, audio and screenshots





The call accepts more; the architecture does not change

The gap between open-weight and frontier keeps narrowing




Not what is possible — who can run it, and where

Shared tool protocols remove the custom glue, not the model




The one shape that is agent news, not model news

Cost per token falls while the scores stay flat





Nothing got smarter; something got affordable

Four of the five never left the LLM region




Only shared tool-use is agent news

Smarter and cheaper are different headlines, so answer them differently




confuse them and you chase a price cut, or miss a real jump

One reopens a decision; the other just changes the price




Reasoning models are not agents

Thinking longer is still one call in, one answer out



still the LLM region




an architecture, not a model

No tool, no second turn, no handing control back



Why this triage habit outlasts any one release


New names arrive; the five shapes do not change



Learn the shapes, not this month's leaderboard

MODULE 11 TAKEAWAYS

You can place any new release on a map you already built



- 1 which region of the map?
- 2 what kind of impact on you?
- 3 five shapes, endless new names

Two questions. Five shapes. One map.

Now open **Module 11 · Lab.** · **Gourav Shah** · **School of DevOps & AI**